

Question 1.1

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import scipy
import scipy.stats as stats
import seaborn as sns
from tqdm import tqdm
```

```
In [2]: df = pd.read_csv('./A1_co2.txt', sep=' ')
df.head()
```

```
Out[2]:
```

	year	month	time	co2
0	1958	3	1958.208	315.71
1	1958	4	1958.292	317.45
2	1958	5	1958.375	317.50
3	1958	6	1958.458	317.10
4	1958	7	1958.542	315.86

You should not use the observations for years 2018 and 2019 (Last 20 observations) for estimations/training - only for comparisons/testing.

```
In [3]: # Let the train set be the measurements before 2018, and the test set be
train = df[df['year'] < 2018]
test = df[df['year'] >= 2018]
```

```
In [4]: # Plot the data
fig, ax = plt.subplots(dpi=800)

# Set the style
sns.set_context("paper", rc={"lines.linewidth": 0.8})
sns.set_palette("colorblind")

ax.scatter(train["time"], train["co2"], label="Training data", s=0.5, c="r")
ax.scatter(test["time"], test["co2"], label="Test data", s=0.5, c="g")


ax.vlines(2018, train["co2"].min(), 425, linestyles="dashed", colors="k")

ax.legend()

ax.set_xlabel("Time [year]")
ax.set_ylabel("$CO_2$ concentration [ppm]")

fig.autofmt_xdate()
fig.show()
```

```
/var/folders/p/_6wwjz4v11fs26vb5j8zwskgm0000gn/T/ipykernel_16591/91593522  
5.py:19: UserWarning: Matplotlib is currently using module://matplotlib_i  
nline.backend_inline, which is a non-GUI backend, so cannot show the figu  
re.  
fig.show()
```



Question 1.2

Question 1.2.1

```
In [5]: MONTHS_IN_YEAR = 12
```

```
x_train = np.vstack([
    [
        np.ones_like(train.index),
        train.index,
        np.sin(2 * np.pi / MONTHS_IN_YEAR * train.index),
        np.cos(2 * np.pi / MONTHS_IN_YEAR * train.index),
    ]
]).T

y_train = train["co2"].values

x_test = np.vstack([
    [
        np.ones_like(test.index),
        test.index,
        np.sin(2 * np.pi / MONTHS_IN_YEAR * test.index),
        np.cos(2 * np.pi / MONTHS_IN_YEAR * test.index),
    ]
]).T

y_test = test["co2"].values
```

```
In [6]: beta_ols, _, _, _ = np.linalg.lstsq(x_train, y_train)
beta_ols
```

```
/var/folders/p/_6wwjz4v11fs26vb5j8zwskgm0000gn/T/ipykernel_16591/2496290795.py:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)`` where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old, explicitly pass `rcond=-1`.
    beta_ols, _, _, _ = np.linalg.lstsq(x_train, y_train)
array([3.06928362e+02, 1.28374044e-01, 1.69090265e+00, 2.25306869e+00])
```

```
Out[6]:
```

Question 1.2.2

```
In [7]: N, p = x_train.shape
```

```
residuals = y_train - x_train @ beta_ols

sigma_ols2 = np.sum(residuals ** 2) / (N - p)
sigma_ols2
```

```
Out[7]: 12.185026694430931
```

```
In [8]: sigma_ols = np.sqrt(sigma_ols2)
sigma_ols
```

```
Out[8]: 3.490705758787316
```


```
In [9]: var_beta_ols = sigma_ols2 * np.linalg.inv(x_train.T @ x_train)
var_beta_ols
```

```
Out[9]: array([[ 6.77484239e-02, -1.41638461e-04, -4.00049374e-04,
   -2.70584607e-04],
 [-1.41638461e-04,  3.95086363e-07,  9.35333307e-07,
  9.35333307e-07],
 [-4.00049374e-04,  9.35333307e-07,  3.39442373e-02,
 -7.99112497e-05],
 [-2.70584607e-04,  9.35333307e-07, -7.99112497e-05,
  3.39442373e-02]])
```

```
In [10]: # Heatmap of the variance of the coefficients
sns.heatmap(np.sqrt(var_beta_ols), annot=True, fmt=".2f")
```

```
/var/folders/p/_6wwjz4v11fs26vb5j8zwskgm0000gn/T/ipykernel_16591/36285869
60.py:2: RuntimeWarning: invalid value encountered in sqrt
    sns.heatmap(np.sqrt(var_beta_ols), annot=True, fmt=".2f")
```

```
Out[10]: <AxesSubplot: >
```



```
In [11]: np.sqrt(np.diag(var_beta_ols))
```

```
Out[11]: array([0.26028527, 0.00062856, 0.18423962, 0.18423962])
```

Question 1.2.3

Question 1.2.4

Plan-of-attack: Relaxtion algorithm

```
In [12]: def rho_matrix(rho: float, n: int):
    """
    Returns the covariance matrix for the observations of a stationary AR
    """
    rhos = np.vander([rho], n, increasing=True)
    return scipy.linalg.toeplitz(rhos)

N, p = X_train.shape

# Initial guess of correlation structure
Sigma_wls = np.eye(N)

# Initial guess of coefficients
beta_ols = beta_wls

betas_wls = []

for _ in range(6):
    # E-step
    Sigma_inv = np.linalg.inv(Sigma_wls)

    H = np.linalg.inv(X_train.T @ Sigma_inv @ X_train) @ X_train.T @ Sigma_inv

    beta_new = H @ y_train

    # M-step
    residuals = y_train - X_train @ beta_new

    rho = np.corrcoef(residuals[:-1], residuals[1:])[0, 1] # 1-lag autocorrelation
    print(rho)

    Sigma_new = rho_matrix(rho, N)

    #if np.allclose(beta, beta_new, atol=1e-6) and np.allclose(Sigma, Sigma_new):
    #    break

    beta_wls = beta_new
    Sigma_wls = Sigma_new

    betas_wls.append(beta_wls)

beta_wls
```

0.9820921007263159
0.9822410526176385
0.9822427429486106
0.982242762307222
0.9822427625289551
0.9822427625314976
Out[12]: array([3.07470864e+02, 1.29839150e-01, 1.66890379e+00, 2.29457132e+00])

Question 1.2.6

```
In [13]: # Calculate the variance of the coefficients
sigma_wls2 = (residuals.T @ np.linalg.inv(Sigma_wls) @ residuals) / (N - 1)
var_beta_wls = np.linalg.inv(X_train.T @ np.linalg.inv(Sigma_wls) @ X_train)
var_beta_wls
```

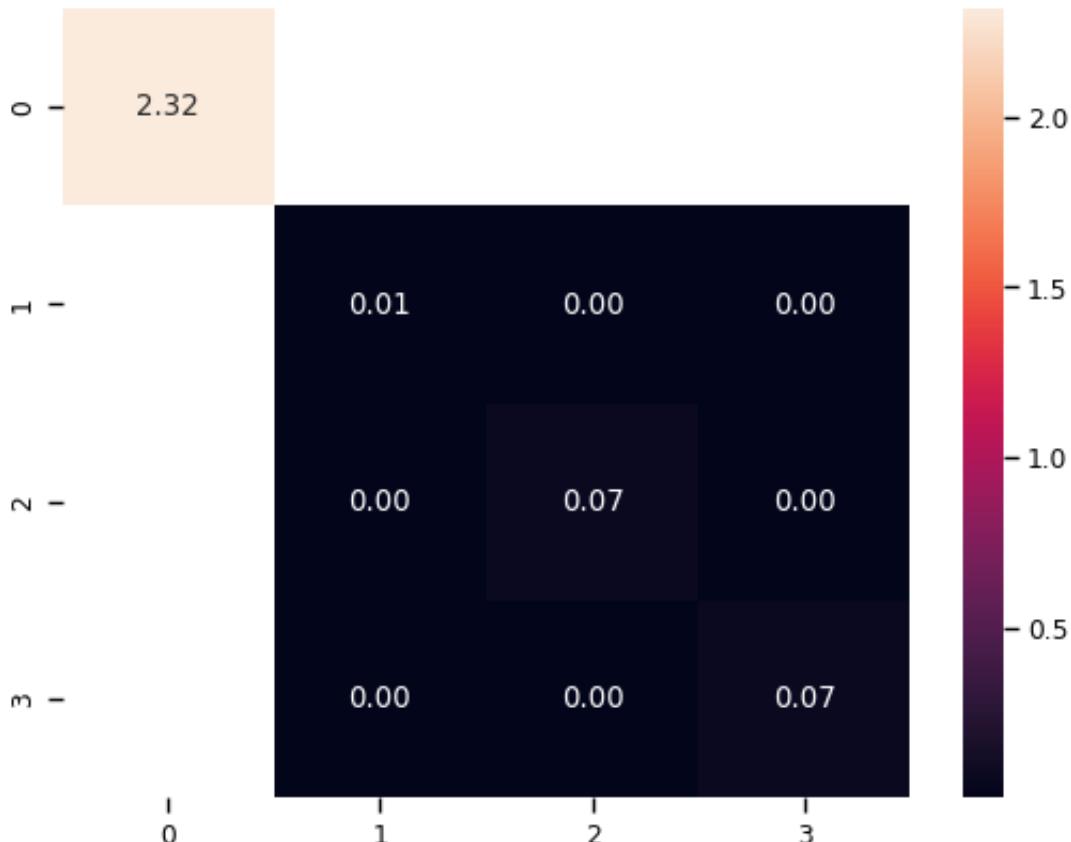
```
Out[13]: array([[ 5.37440920e+00, -1.03811751e-02, -5.30015696e-04,
                   -1.12349504e-03],
                  [-1.03811751e-02,  2.89572528e-05,  2.30615165e-06,
                   2.30615165e-06],
                  [-5.30015696e-04,  2.30615165e-06,  4.57062027e-03,
                   1.28481766e-05],
                  [-1.12349504e-03,  2.30615165e-06,  1.28481766e-05,
                   4.57062027e-03]])
```

```
In [14]: sigma_wls2
```

```
Out[14]: 12.268130156533738
```

```
In [15]: sns.heatmap(np.sqrt(var_beta_wls), annot=True, fmt=".2f")
```

```
/var/folders/p_/_6wwjz4v11fs26vb5j8zwskgm0000gn/T/ipykernel_16591/1372562171.py:1: RuntimeWarning: invalid value encountered in sqrt
    sns.heatmap(np.sqrt(var_beta_wls), annot=True, fmt=".2f")
<AxesSubplot: >
```



```
In [16]: np.sqrt(np.diag(var_beta_wls))
```

```
Out[16]: array([2.3182772 , 0.00538119, 0.06760636, 0.06760636])
```

```
In [17]: beta_ols - beta_wls
```

```
Out[17]: array([-0.54250225, -0.00146511,  0.02199886, -0.04150264])
```

Model comparison plot

```
In [18]:
```

```
y_hat_ols = X_train @ beta_ols
y_hat_wls = X_train @ beta_wls

y_test_hat_ols = X_test @ beta_ols
y_test_hat_wls = X_test @ beta_wls

# Plot the train set and the fitted curves
fig, ax = plt.subplots(dpi=800)

# Set the style
sns.set_context("paper", rc={"lines.linewidth": 0.8})
sns.set_palette("colorblind")

ax.scatter(train["time"], train["co2"], label="Training data", s=0.5, c="r")
ax.scatter(test["time"], test["co2"], label="Test data", s=0.5, c="g")

ax.plot(train["time"], y_hat_ols, label="OLS fit", alpha=0.8, c="C0")
ax.plot(test["time"], y_test_hat_ols, alpha=0.8, c="C0")

ax.plot(train["time"], y_hat_wls, label="WLS fit", alpha=0.8, c="C1")
ax.plot(test["time"], y_test_hat_wls, alpha=0.8, c="C1")


ax.vlines(2018, 300, 420, linestyles="dashed", label="Train/Test split",
          ax.legend()

          ax.set_xlabel("Time [year]")
          ax.set_ylabel("$CO_2$ concentration [ppm]")

          fig.autofmt_xdate()
          fig.show()
```

```
/var/folders/p/_6wwjz4v11fs26vb5j8zwskgm0000gn/T/ipykernel_16591/3738848100.py:31: UserWarning: Matplotlib is currently using module://matplotlib_inline.backend_inline, which is a non-GUI backend, so cannot show the figure.
```

```
    fig.show()
```



```
In [19]: y_hat_ols = X_train @ beta_ols
y_hat_wls = X_train @ beta_wls

# Plot the train set and the fitted curves
fig, ax = plt.subplots(dpi=800)

# Set the style
sns.set_context("paper", rc={"lines.linewidth": 0.8})
sns.set_palette("colorblind")

idx_2010 = train["year"] > 2010

ax.scatter(train["time"][idx_2010], train["co2"][idx_2010], label="Training data")
ax.scatter(test["time"], test["co2"], label="Test data", s=0.5, c="g")

ax.plot(train["time"][idx_2010], y_hat_ols[idx_2010], label="OLS fit", alpha=0.8, c="C0")
ax.plot(test["time"], y_hat_ols, alpha=0.8, c="C0")

ax.plot(train["time"][idx_2010], y_hat_wls[idx_2010], label="WLS fit", alpha=0.8, c="C1")
ax.plot(test["time"], y_hat_wls, alpha=0.8, c="C1")


ax.vlines(test["time"].min(), 380, 420, linestyles="dashed", label="Test set split")

ax.legend()

ax.set_xlabel("Time [year]")
ax.set_ylabel("$CO_2$ concentration [ppm]")

fig.autofmt_xdate()
fig.show()
```

```
/var/folders/p/_6wwjz4v11fs26vb5j8zwskgm0000gn/T/ipykernel_16591/39755760  
35.py:31: UserWarning: Matplotlib is currently using module://matplotlib_  
inline.backend_inline, which is a non-GUI backend, so cannot show the fig-  
ure.  
fig.show()
```



Model comparison

```
In [20]: residuals_ols = train["co2"] - y_hat_ols  
residuals_wls = train["co2"] - y_hat_wls
```

```
In [21]: # Scatter plot of the residuals for the OLS and WLS fits
fig, ax = plt.subplots(ncols=2, sharey=True, dpi=800)

# Set the style
sns.set_context("paper", rc={"lines.linewidth": 0.8})
sns.set_palette("colorblind")

ax[0].scatter(train["time"], residuals_ols, label="OLS residuals", s=0.5,
ax[1].scatter(train["time"], residuals_wls, label="WLS residuals", s=0.5)

ax[0].hlines(0, train["time"].min(), train["time"].max(), linestyle="--",
ax[1].hlines(0, train["time"].min(), train["time"].max(), linestyle="--",

ax[0].legend()
ax[1].legend()


ax[0].set_xlabel("Time [year]")
ax[1].set_xlabel("Time [year]")

ax[0].set_ylabel("Residuals")
ax[1].set_ylabel("Residuals")

fig.autofmt_xdate()
fig.show()
```

```
/var/folders/p/_6wwjz4v11fs26vb5j8zwskgm0000gn/T/ipykernel_16591/10314583
36.py:24: UserWarning: Matplotlib is currently using module://matplotlib_
inline.backend_inline, which is a non-GUI backend, so cannot show the fig-
ure.
```

```
    fig.show()
```



```
In [ ]:
```

Question 1.3

```
In [22]: plt.rc("axes", axisbelow=True)
plt.rcParams["figure.dpi"] = 400
plt.rcParams["figure.figsize"] = (9, 4)
sns.set_palette("colorblind")
```

Question 1.3.2

Filter the data with the chosen model.

```
In [23]: # Load data

year, month, time, co2 = df.values.T

year = year.astype(int)
month = month.astype(int)

# Index data by indices. 12 months per year
# NOTE: There is a constant time delta between each data point,
# so indexing directly by indices is fine.
p = 12

# Split into test and train
test_idx = year >= 2018
time_test, co2_test = time[test_idx], co2[test_idx]
time_train, co2_train = time[-test_idx], co2[-test_idx]
y_train = co2_train[:, None]

train_n = len(time_train)
test_n = len(time_test)

param_n = 4
```

```
In [24]: def llt_pred_interval(y_pred, t, F, f, var, alpha=0.05):
    return y_pred + np.array([1, -1]) * stats.t.ppf(
        1 - alpha / 2, t - param_n
    ) * np.sqrt(var) * np.sqrt(1 + f.T @ np.linalg.inv(F) @ f)

def f(t):
    return np.array([
        [
            1,
            t,
            np.sin(2 * np.pi / p * t),
            np.cos(2 * np.pi / p * t),
        ]
    ])
) .T

def llt_predict(start=10, lamb=0.9):
    assert start > 0, "Start must be greater than 0"
```

```

# Create np.linalg.inverse transition matrix
L_inv = np.linalg.inv(
    np.array(
        [
            [1, 0, 0, 0],
            [1, 1, 0, 0],
            [0, 0, np.cos(2 * np.pi / p), np.sin(2 * np.pi / p)],
            [0, 0, -np.sin(2 * np.pi / p), np.cos(2 * np.pi / p)],
        ]
    )
)

# Create design matrix for t'th time step
f_0 = f(0)

F = np.zeros((param_n, param_n))
h = np.zeros((param_n, 1))

# Store predictions and variances
y_pred_train = []
y_pred_train_interval = []
y_pred_train_mean = []
y_pred_train_var = []
y_pred_train_var_sum = 0
thetas = []

# Go through each data point and make one-step predictions
for i in range(0, train_n):
    # If burn-in period is over, make predictions
    if i >= start:
        F_inv = np.linalg.inv(F)

        # Calculate parameters for previous time step
        theta = F_inv @ h

        # Make one step prediction using parameters
        # from previous data to predict current data.
        y_pred = f(1).T @ theta

        # Calculate noise variance from next prediction
        y_pred_train_var_sum += (
            (y_train[i, 0] - y_pred[0]) ** 2 / (1 + f(1).T @ F_inv @
            )[0, 0]
        )
        y_pred_train_var.append(y_pred_train_var_sum / (i - start + 1))

        # Store predictions
        y_pred_train.append(y_pred[0, 0])
        y_pred_train_mean.append(theta[0, 0])
        thetas.append(theta)

        # Store prediction interval
        y_pred_train_interval.append(
            llt_pred_interval(y_pred, i - 1, F, f(1), y_pred_train_var
            )
        )

        # Update F and h
        F = F + lamb**i * f(-i) @ f(-i).T
        h = lamb * L_inv @ h + f_0 @ y_train[i : i + 1]

# Calculate final parameters

```

```

theta = np.linalg.inv(F) @ h
thetas.append(theta)

# Make predictions for test data
y_pred_test = np.array([f(t + 1).squeeze() for t in range(test_n)]) @

# Make intervals for test data
y_pred_test_interval = np.array(
    [
        llt_pred_interval(
            y_pred_test[t], train_n, F, f(t + 1), y_pred_train_var[-1]
        )
        for t in range(test_n)
    ]
).squeeze()

# Return a big mess
return (
    np.array(y_pred_train),
    np.array(y_pred_train_interval).squeeze(),
    np.array(y_pred_train_mean),
    np.array(y_pred_train_var),
    y_pred_test,
    y_pred_test_interval,
)
)

# Amount of data to skip predictions for
burn_in_1 = 10

# Retrieve predictions
(
    llt_y_pred_train,
    llt_y_pred_train_interval,
    llt_y_pred_train_mean,
    llt_y_pred_train_var,
    llt_y_pred_test,
    llt_y_pred_test_interval,
) = llt_predict(start=burn_in_1, lamb=0.9)

```

```
In [25]: llt_residuals = y_train[burn_in_1:, 0] - llt_y_pred_train

# Plot residuals as function of y
plt.scatter(y_train[burn_in_1:, 0], llt_residuals, s=0.8, c='r', label='R')

plt.xlabel("CO2 concentration [ppm]")
plt.ylabel("Residuals [ppm]")

plt.legend()
plt.grid()

plt.show()

# Plot residuals as function of time
plt.plot(time_train[burn_in_1:], llt_residuals, '--', linewidth=0.4)
plt.scatter(time_train[burn_in_1:], llt_residuals, c='r', s=0.8, label='R')

plt.xlabel("Time [year]")
plt.ylabel("Residuals [ppm]")

plt.legend()
plt.grid()

plt.show()

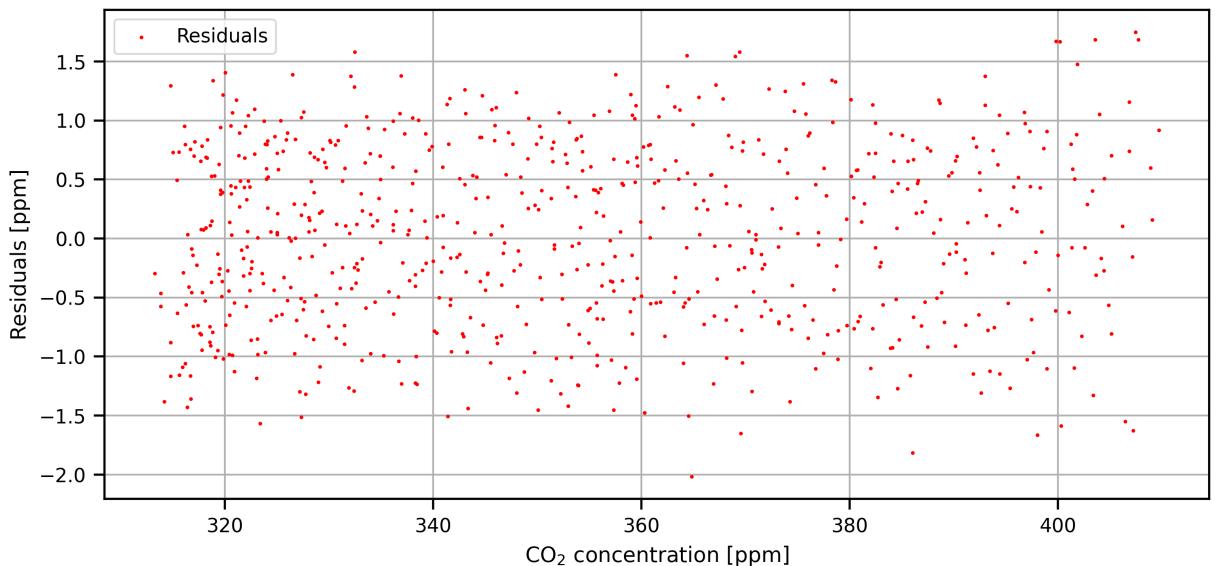
# Plot variance
plt.scatter(time_train[burn_in_1:], np.sqrt(llt_y_pred_train_var), s=0.8, c='r', label='S')

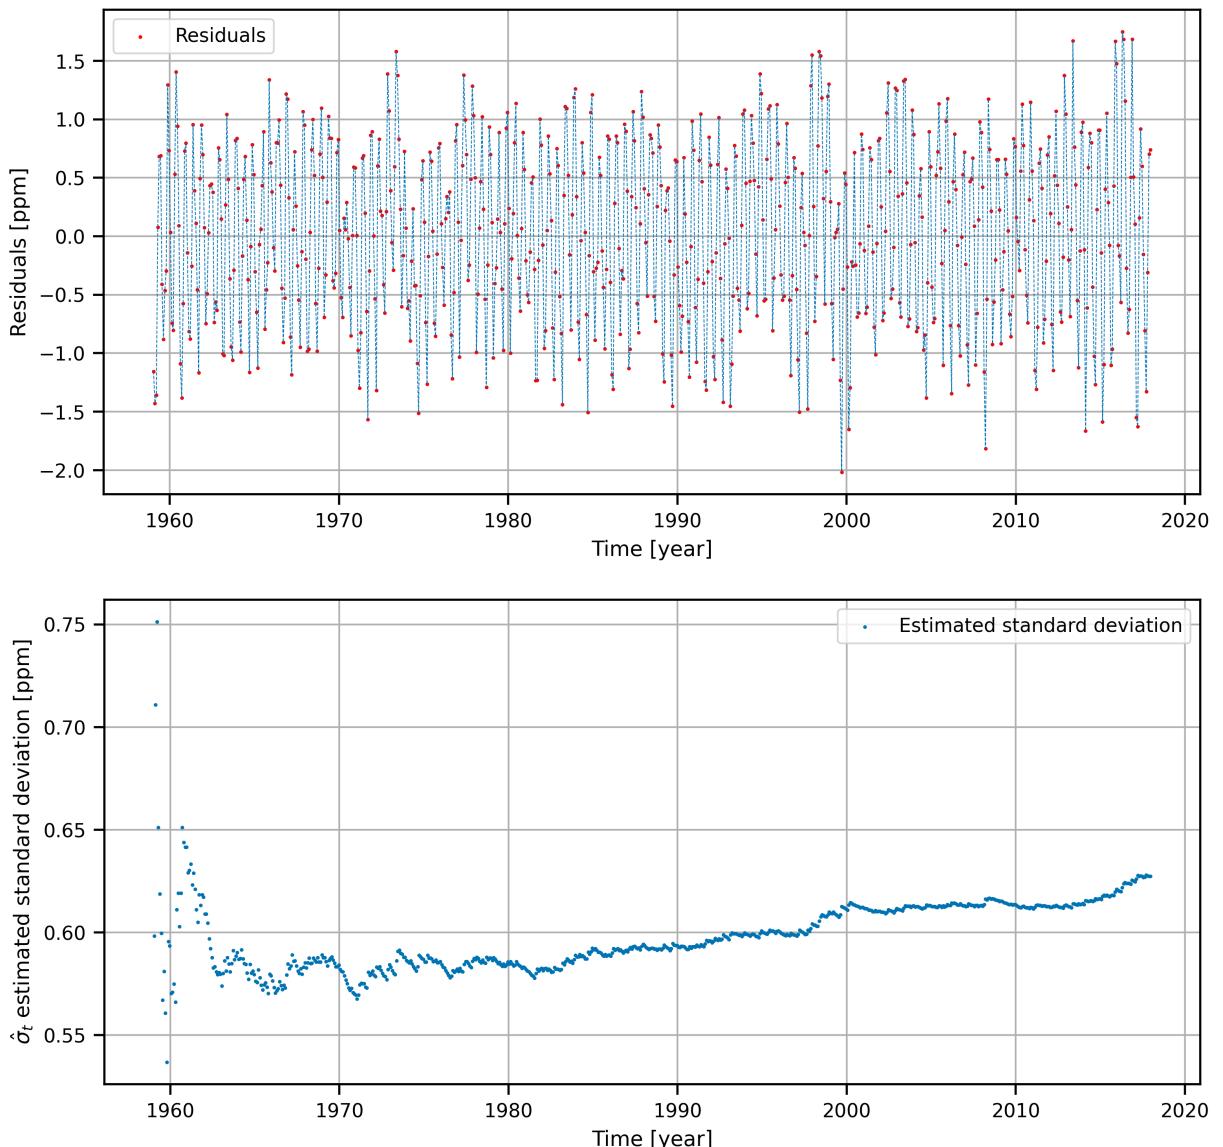
plt.xlabel("Time [year]")
plt.ylabel("$\hat{\sigma}_t$ estimated standard deviation [ppm]")

plt.legend()
plt.grid()

plt.show()

print(np.sqrt(llt_y_pred_train_var[-1]))
```





0.6273878545097442

1.4

```
In [26]: # Set ranges of points to skip and lambdas to test
burn_in_2 = 100
lambs = np.linspace(0+2e-4, 1-9e-3, 1000)

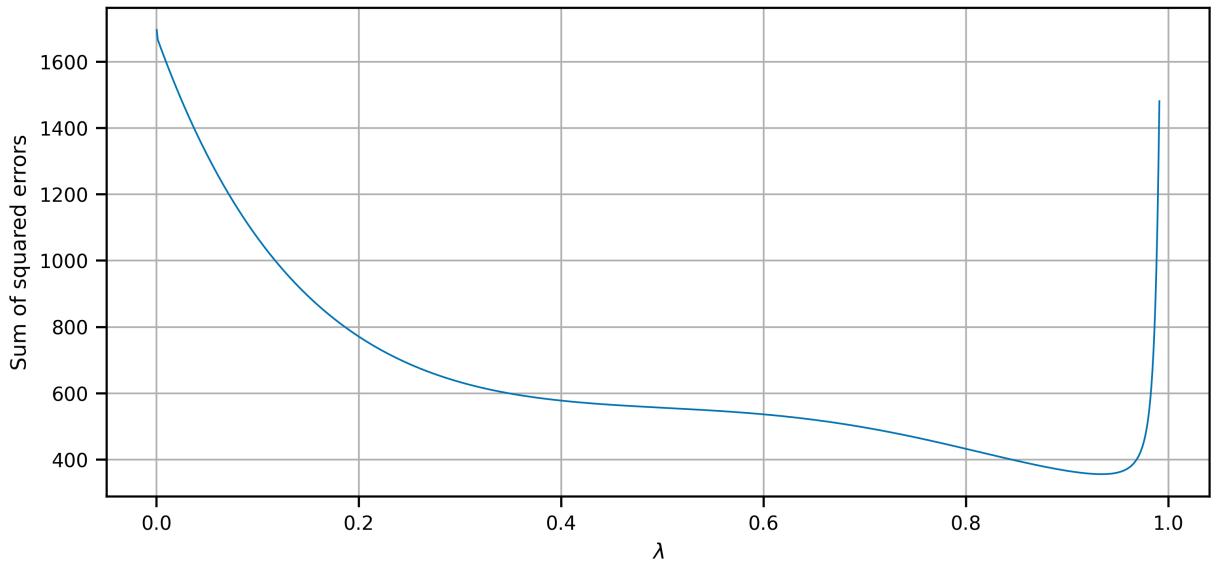
# Predict and calculate squared error for each lambda
sqerrors = [((y_train[burn_in_2:, 0] - llt_predict(start=burn_in_2, lambd=
# Plot squared error as a function of lambda
plt.plot(lambs, sqerrors)

plt.grid()
plt.xlabel('$\lambda$')
plt.ylabel('Sum of squared errors')

plt.show()

print(lambs[800+np.argmin(sqerrors[800:])]))

100% |██████████| 1000/1000 [01:00<00:00, 16.53it/s]
```



0.933476076076076

```
In [27]: optimal_lambda = lambs[800 + np.argmin(sqerrors[800:])]
print(optimal_lambda)
print(sqerrors[800 + np.argmin(sqerrors[800:]):])

fig, ax = plt.subplots(figsize=(10, 5))
ax.grid()

# Plot squared error as a function of lambda
ax.plot(lambs, sqerrors, label="Sum of squared error")
ax.vlines(
    optimal_lambda, 200, 900, color="r", linestyles="dashed", label="Optimal")
    
```

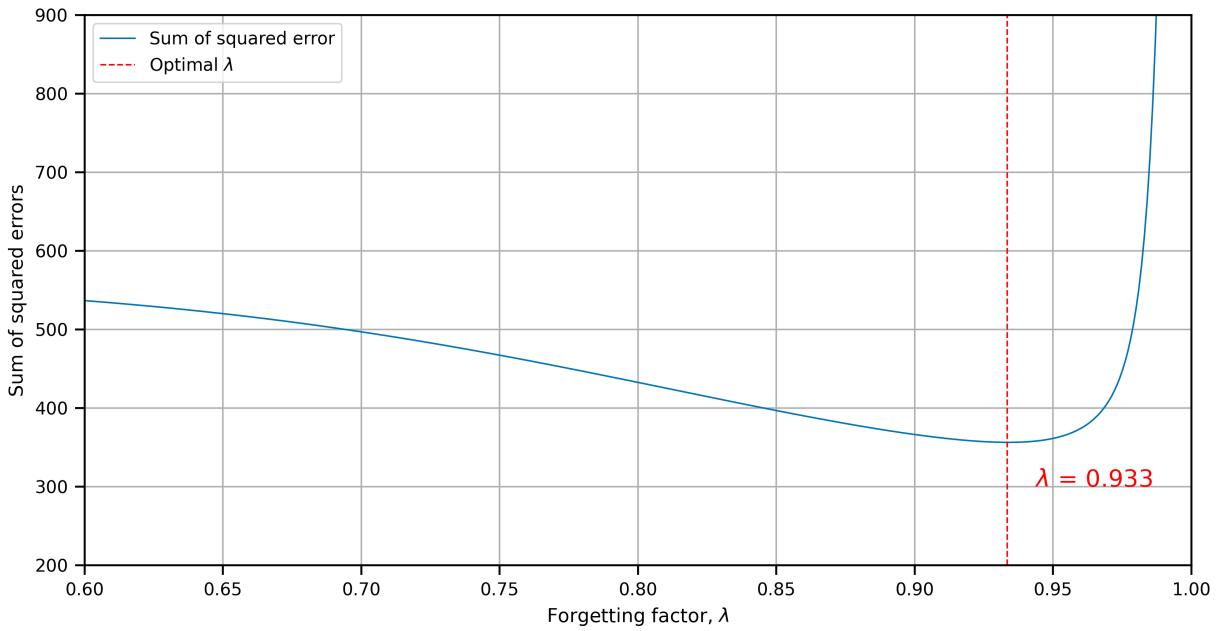
plt.text(
 optimal_lambda + 0.01,
 300,
 f"\$\lambda\$ = {optimal_lambda:.3f}",
 color="r",
 fontsize=12,
)

ax.set_xlim(0.6, 1)
ax.set_ylim(200, 900)

ax.set_xlabel("Forgetting factor, \$\lambda\$")
ax.set_ylabel("Sum of squared errors")
ax.legend()

plt.show()

0.933476076076076
356.2038625197131



1.1

```
In [28]: # Set plot options
plot_model_params = {"alpha": 0.7, "linestyle": '-', "linewidth": 0.8}
plot_interval_params = {"alpha": 0.5, "linestyle": '--', "linewidth": 0.4}
plot_data_params = {"alpha": 0.9, "s": 0.2}

# Plot data
plt.scatter(time_train, co2_train, label="Training", c="r", **plot_data_p
plt.scatter(np.concatenate([time_train[-1], None], time_test)), np.concatenate([co2_train[-1], None], co2_test), label="Test", c="b", **plot_data_p)

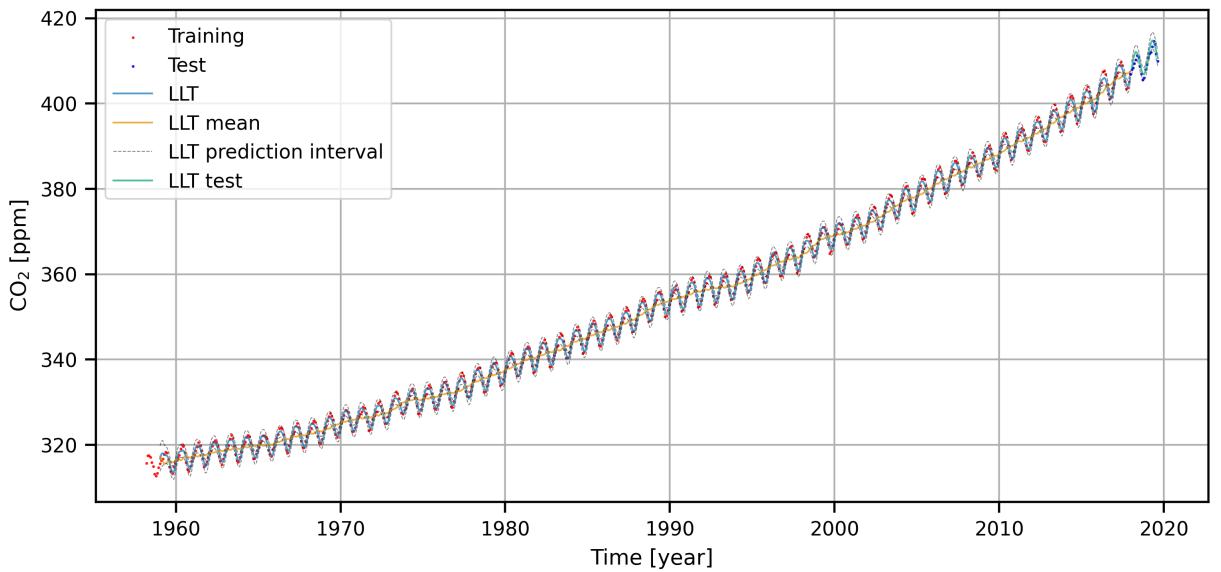
# Plot model predictions
plt.plot(time_train[burn_in_1:], llt_y_pred_train, label="LLT", **plot_mo
plt.plot(time_train[burn_in_1:], llt_y_pred_train_mean, label="LLT mean", c='k', **plot_mo
plt.plot(time_train[burn_in_1:], llt_y_pred_train_interval[:, 0], label="LLT interval", **plot_mo
plt.plot(time_train[burn_in_1:], llt_y_pred_train_interval[:, 1], c='k', **plot_mo

plt.plot(time_test, llt_y_pred_test, label="LLT test", **plot_model_params)
plt.plot(time_test, llt_y_pred_test_interval[:, 0], c='k', **plot_interval_params)
plt.plot(time_test, llt_y_pred_test_interval[:, 1], c='k', **plot_interval_params)

# Plot
plt.xlabel("Time [year]")
plt.ylabel("CO$ 2$ [ppm]")

plt.legend()
plt.grid()

plt.show()
```



```
In [29]: zoom_idx = np.argwhere(year[-test_idx] >= 2009).squeeze()

# Plot data
plt.scatter(time_train[zoom_idx], co2_train[zoom_idx], label="Training",
            c='red')
plt.scatter(np.concatenate([time_train[-1], None], time_test)), np.concatenate(
    llt_y_pred_train[zoom_idx], llt_y_pred_test)

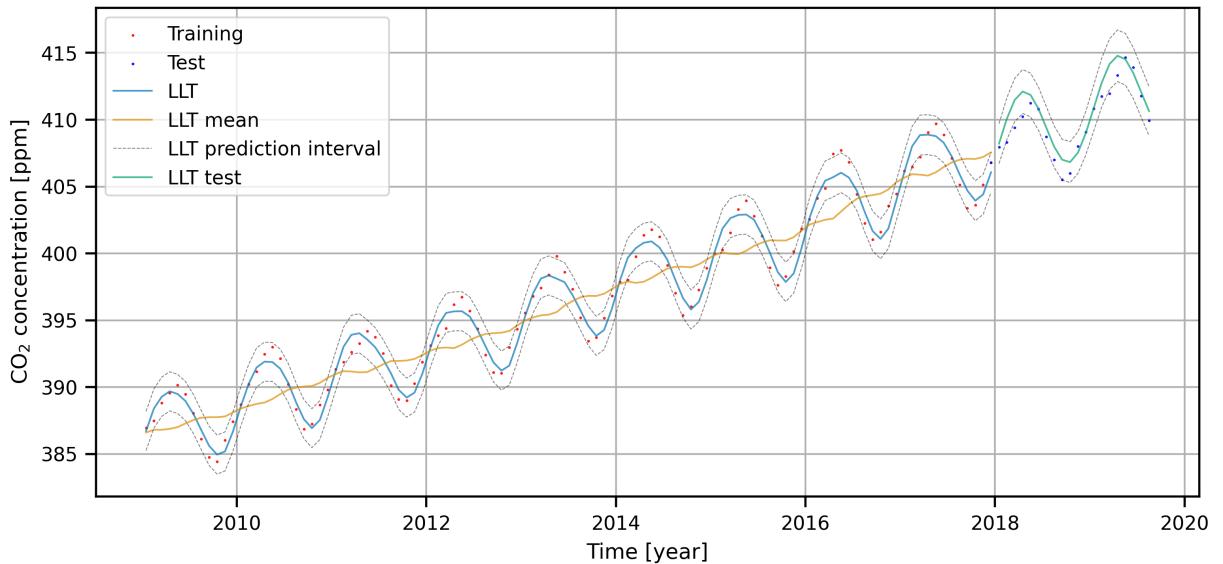
# Plot model predictions
plt.plot(time_train[zoom_idx], llt_y_pred_train[zoom_idx-burn_in_1], label="LLT",
          c='blue')
plt.plot(time_train[zoom_idx], llt_y_pred_train_mean[zoom_idx-burn_in_1], label="LLT mean",
          c='orange')
plt.plot(time_train[zoom_idx], llt_y_pred_train_interval[zoom_idx-burn_in_1], label="LLT prediction interval",
          c='grey', ls='dashed')
plt.plot(time_train[zoom_idx], llt_y_pred_train_interval[zoom_idx-burn_in_1], label="LLT prediction interval",
          c='grey', ls='dashed')

plt.plot(time_test, llt_y_pred_test, label="LLT test", **plot_model_params)
plt.plot(time_test, llt_y_pred_test_interval[:, 0], c='k', **plot_interval_params)
plt.plot(time_test, llt_y_pred_test_interval[:, 1], c='k', **plot_interval_params)

# Plot
plt.xlabel("Time [year]")
plt.ylabel("CO2 concentration [ppm]")

plt.legend()
plt.grid()

plt.show()
```



```
In [30]: # Print table of future predictions
y_pred_test_idx = np.array([1, 2, 6, 12, 20])-1

for t, true, pred, lower, upper in zip(y_pred_test_idx+1, co2_test[y_pred
    print(f"{t} & {true:.2f} & {pred:.2f} & {lower:.2f} & {upper:.2f}\\\\\\
1 & 407.96 & 408.21 & 406.72 & 409.71\\\
2 & 408.32 & 410.08 & 408.51 & 411.64\\\
6 & 410.79 & 410.83 & 409.22 & 412.44\\\
12 & 409.07 & 409.00 & 407.37 & 410.62\\\
20 & 409.95 & 410.64 & 408.81 & 412.48\\\

```

Question 1.4.2

```
In [31]: BURN_IN = 10
optimal_llt_y_pred_train, optimal_llt_y_pred_train_interval, optimal_llt_

# Set plot options
plot_model_params = {"alpha": 0.7, "linestyle": '-', "linewidth": 0.8}
plot_interval_params = {"alpha": 0.5, "linestyle": '--', "linewidth": 0.4}
plot_data_params = {"alpha": 0.9, "s": 0.2}

# Plot data
plt.scatter(time_train, co2_train, label="Training", c="r", **plot_data_p
plt.scatter(np.concatenate([time_train[-1], None], time_test)), np.concatenate


# Plot model predictions
plt.plot(time_train[BURN_IN:], optimal_llt_y_pred_train, label="LLT", **p
plt.plot(time_train[BURN_IN:], optimal_llt_y_pred_train_mean, label="LLT
plt.plot(time_train[BURN_IN:], optimal_llt_y_pred_train_interval[:, 0], l
plt.plot(time_train[BURN_IN:], optimal_llt_y_pred_train_interval[:, 1], c

plt.plot(time_test, optimal_llt_y_pred_test, label="LLT test", **plot_mod
plt.plot(time_test, optimal_llt_y_pred_test_interval[:, 0], c='k', **plot
plt.plot(time_test, optimal_llt_y_pred_test_interval[:, 1], c='k', **plot

# Plot
plt.xlabel("Time [year]")
plt.ylabel("CO2 [ppm]")

plt.legend()
plt.grid()

plt.show()
```



```
In [32]: zoom_idx = np.argwhere(year[-test_idx] >= 2009).squeeze()

# Plot data
plt.scatter(time_train[zoom_idx], co2_train[zoom_idx], label="Training",
            plt.scatter(np.concatenate([time_train[-1, None], time_test]), np.concatenate([co2_train[zoom_idx], co2_test]), label="Test")

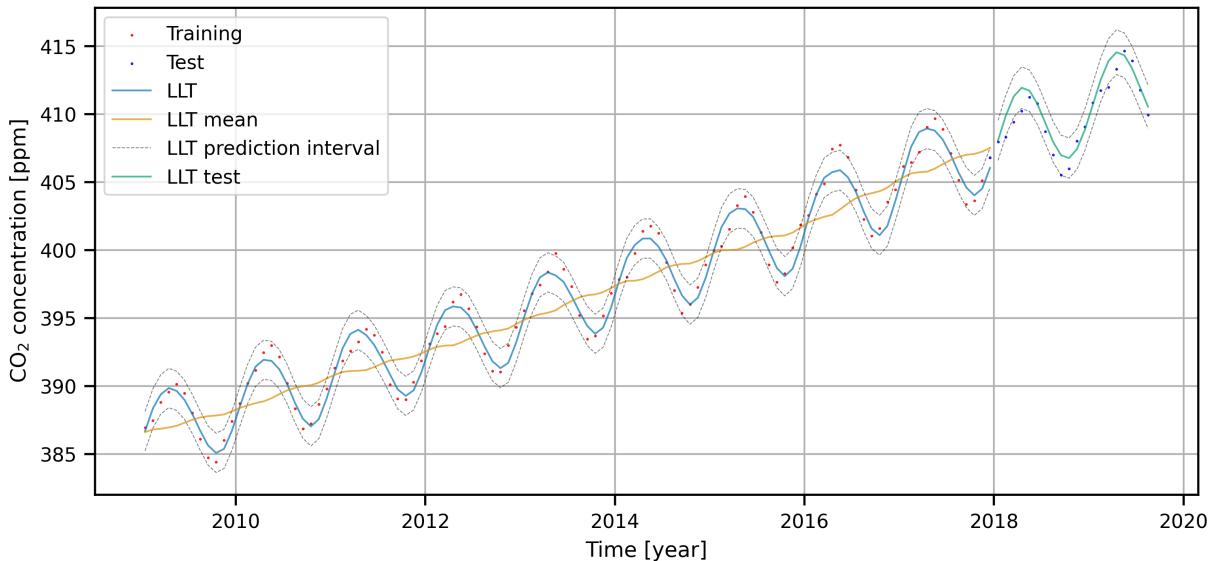
# Plot model predictions
plt.plot(time_train[zoom_idx], optimal_llt_y_pred_train[zoom_idx-BURN_IN:], label="LLT")
plt.plot(time_train[zoom_idx], optimal_llt_y_pred_train_mean[zoom_idx-BURN_IN:], label="LLT mean")
plt.plot(time_train[zoom_idx], optimal_llt_y_pred_train_interval[zoom_idx-BURN_IN:], label="LLT prediction interval")
plt.plot(time_train[zoom_idx], optimal_llt_y_pred_train_interval[zoom_idx-BURN_IN:], label="LLT test")

plt.plot(time_test, optimal_llt_y_pred_test, label="LLT test", **plot_mod)
plt.plot(time_test, optimal_llt_y_pred_test_interval[:, 0], c='k', **plot_mod)
plt.plot(time_test, optimal_llt_y_pred_test_interval[:, 1], c='k', **plot_mod)

# Plot
plt.xlabel("Time [year]")
plt.ylabel("CO2 concentration [ppm]")

plt.legend()
plt.grid()

plt.show()
```



```
In [33]: # Print table of future predictions
y_pred_test_idx = np.array([1, 2, 6, 12, 20])-1

for t, true, pred, lower, upper in zip(y_pred_test_idx+1, co2_test[y_pred_test_idx], optimal_llt_y_pred_test[y_pred_test_idx], optimal_llt_y_pred_test_interval[y_pred_test_idx, 0], optimal_llt_y_pred_test_interval[y_pred_test_idx, 1]):
    print(f"{t} & {true:.2f} & {pred:.2f} & {lower:.2f} & {upper:.2f} \\ \\ \\"
1 & 407.96 & 408.07 & 406.60 & 409.54 \\
2 & 408.32 & 409.90 & 408.40 & 411.41 \\
6 & 410.79 & 410.74 & 409.23 & 412.25 \\
12 & 409.07 & 408.84 & 407.32 & 410.36 \\
20 & 409.95 & 410.54 & 408.95 & 412.14 \\ \\ \\
```