

Time Series Analysis - Assignment 4

Lauge Thode Hermansen
s204111

Tymoteusz Barcinski
s221937

Vincent Olesen
s184017

May 2023

1 Introduction

This report aims to model and filter water salinity for Roskilde Fjord using state-space models.

The data used in the report was collected from sensors on a bouy located in the water near Kulhuse - in the opening of Roskilde Fjord. The measurements consist of water salinity and dissolved oxygen, which we will focus on, and seven other sensor measurements. The data was collected from 2017-08-24 to 2017-12-06 with a sampling rate of 30 minutes. The data has missing values; notably, there is a missing observation period from 2017-09-26 at 12:00:00 to 2017-09-28 at 16:00:00, but other missing observations exist, notably in September and at the end of November.




Figure 1: Plot of salinity for the whole period (left) and the last week (right)




Figure 2: Plot of salinity for the whole period but with the y-range set to [15,22] for better viewing of the variation, The red lines, represent the time index of the observations not shown on the plot, because they are too low.




Figure 3: Plot of dissolved oxygen for the whole period (left) and the last week (right)




Figure 4: Plot of dissolved oxygen for the whole period

In figs. 1 and 3, we have the entire period (left) and the last week (right) to highlight daily trends.

Considering the entire period, there are outliers, notably in September and at the end of November.

We see a higher variation of the dissolved oxygen before the end of October (the period of the missing values) compared to the rest of the period, possibly due to a sensor replacement. We see a linear trend in the measured dissolved oxygen from the end of October (the period of the missing values).

Interestingly, the salinity has lower variance before the missing period, and a higher variance after the sensor replacement.

2 Random walk state-space model of salinity

We want to model the salinity X_t at time-step t , as a random walk process that is observed at equidistant time points. We take a similar approach as in [1, Ex. 10.6]. Our model for the state is

$$X_t = X_{t-1} + \varepsilon_{1,t}$$

And when we model our observations as just the state with some measurement noise, we get the state-space model

$$\begin{cases} \mathbf{X}_t = \mathbf{A}\mathbf{X}_{t-1} + \varepsilon_{1,t} & (\text{No input}) \\ \mathbf{Y}_t = \mathbf{C}\mathbf{X}_t + \varepsilon_{2,t} \\ \mathbf{A} = \begin{pmatrix} 1 \end{pmatrix} \\ \mathbf{C} = \begin{pmatrix} 1 \end{pmatrix} \\ \Sigma_1 = \begin{pmatrix} \sigma_1^2 \end{pmatrix} \\ \Sigma_2 = \begin{pmatrix} \sigma_2^2 \end{pmatrix} \end{cases} \quad (1)$$

Where σ_1^2 is the system variance and σ_2^2 is the observation variance. In order to compute prediction intervals later on, we assume that the system and measurement errors are normally distributed with mean zero.

Note that the matrices are time-invariant and therefore the subscripts have been omitted. As C has full rank, the system is observable.

3 Pure Kalman filter




Figure 5: Predictions of the Pure Kalman Filter along with the prediction intervals and standardized one-step predictions errors. The blue points are points we deem to be outliers, and have been marked such that their impact upon the Kalman filter can be more readily be seen.




Figure 6: Predictions of the Pure Kalman Filter along with the prediction intervals and standardized one-step predictions errors in the range $[800, 950]$. The blue points are points we deem to be outliers, and have been marked such that their impact upon the Kalman filter can be more readily be seen.

When looking at the plots we see that the predictions follow the observations quite well. The empirical coverage (probability of being inside the confidence interval) is 95.52% which is actually pretty close to the desired 95%.

Considering the standardized errors, we see that the outliers seem to throw off the filter. Usually it takes 3-4 observations until the predictions are reasonable and the variance estimates are well calibrated again.

Upon close examination of the interval from 800 to 950 in fig. 6 we see that the prediction (the green line) has a tendency to be just the previous observation. This can be explained by looking at the prediction step in the Kalman filter, where we simplify the equations, as we don't have any control input.

Predicting the state, X_{t+1} , is based on the reconstructed state.

$$\hat{X}_{t+1|t} = A\hat{X}_{t|t} \quad (2)$$

Predicting the observation Y_{t+1} is done by

$$\hat{Y}_{t+1|t} = C\hat{X}_{t+1|t} \quad (3)$$

Since in this case, $A = C = 1$, we get

$$\hat{Y}_{t+1|t} = C\hat{X}_{t+1|t} = \hat{X}_{t|t} \quad (4)$$

So the prediction of the next observation is actually just the reconstruction of the previous state. Therefore, if our reconstructions are good estimates of the true states, the results in the plots are actually exactly what we would expect.

Empirically, the kalman gain, K , quickly converges to $K \approx 0.7320508$, when there are no missing values, which means that

$$\begin{aligned} \hat{Y}_{t+1|t} &= \hat{X}_{t|t} \\ &= \hat{X}_{t|t-1} + K_t(Y_t - C\hat{X}_{t|t-1}) \\ &= \hat{X}_{t|t-1} + 0.7320508(Y_t - 1\hat{X}_{t|t-1}) \\ &= 0.2679492\hat{X}_{t|t-1} + 0.7320508Y_t \end{aligned}$$

So the one step prediction is dominated by the previous observation, which aligns well with what we see in the plot. Since the salinity dynamics are modelled as a random walk, using the Kalman filter for prediction does not say much more about the salinity level than the previous observation. The uncertainty of the prediction may, however, be useful as a gauge for how rapidly the system may change.

| State | $\hat{X}_{5000 5000}$ | $\Sigma_{5000 5000}^{xx}$ | $\hat{X}_{5001 5000}$ | $\Sigma_{5001 5000}^{xx}$ | $\Sigma_{5001 5000}^{yy}$ |
|-------|-----------------------|---------------------------|-----------------------|---------------------------|---------------------------|
| Value | 20.76 | $3.660 \cdot 10^{-3}$ | 20.76 | $1.366 \cdot 10^{-2}$ | $1.866 \cdot 10^{-2}$ |

Table 1: Final state of filter at observation 5000, where all observations have been included.

4 Skipping outliers when filtering

Observations are deemed outliers if they are 6 standard deviations away from a prediction¹. Outliers are treated the same as missing observations. Formally, outliers are defined as

$$\mathcal{Y}_{\text{obs}} = \{Y_t\}_{t=1}^{5000}$$

$$\mathcal{Y}_{\text{outlier}} = \left\{ Y_t \in \mathcal{Y}_{\text{obs}} \left| \frac{|Y_t - C\hat{X}_{t|t-1}|}{\Sigma_{t|t-1}^{yy}} > 6 \right. \right\}$$

There are 10 outliers and 111 missing observations, which means in total 121 observations are skipped. The missing observations are usually found in clusters. Skipped observations are handled by setting the associated reconstruction to the same time point's prediction, i.e. if the observation at time t was skipped,

$$\hat{X}_{t|t} = \hat{X}_{t|t-1}$$

$$\Sigma_{t|t}^{xx} = \Sigma_{t|t-1}^{xx}.$$

| t | 814 | 1203 | 1296 | 2733 | 3692 | 4038 | 4578 | 4607 | 4609 | 4686 |
|----------------|-------|-------|------|-------|-------|-------|-------|------|------|-------|
| \mathbf{Y}_t | 15.08 | 15.17 | 0.43 | 18.24 | 17.81 | 17.95 | 17.17 | 1.37 | 1.32 | 18.13 |

Table 2: All outlier indexes and their values.

As expected, the outliers appear to form at time points where the salinity suddenly changes rapidly compared to the variance of the local past. Afterwards, the filter appears to adapt to this increase in variance, such that subsequent nearby rapidly changing salinity values are not marked as outliers.

However, based on the previous plots with clear outliers, the outlier detection may be overzealous as it appears to have perfect recall but mediocre precision, meaning that it classifies as outliers observations that have high variance due to the dynamics of a system, instead of being outliers. The threshold should therefore perhaps be increased from 6.

¹In the multivariate case, if the noise is Gaussian, the Mahalanobis distance can maybe be used to decide whether a prediction is an outlier.

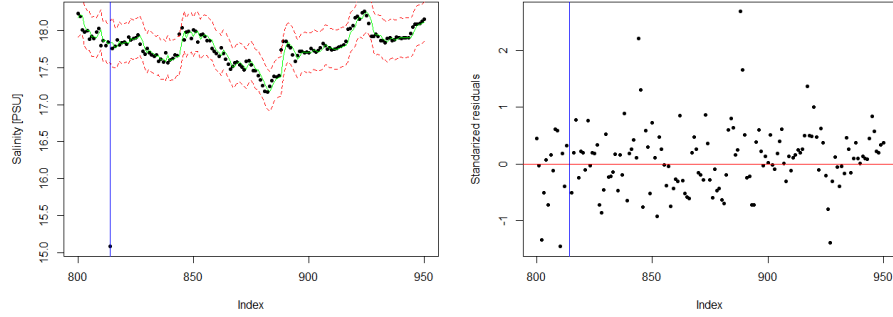


Figure 7: Predictions of the Kalman Filter with the skipping outliers mechanism along with the prediction intervals and standardized one-step prediction errors in the range $[800, 950]$. the blue lines represent the index of the skipped values

Compared to fig. 6, the Kalman Filter is no longer thrown off temporarily by the large spike in variance caused by the outlier. This is evident in the residuals plot. Other than that, things look very similar.




Figure 8: ACF and PACF plots for the standardized one-step prediction residuals for the entire data set.

| State | $\hat{X}_{5000 5000}$ | $\Sigma_{5000 5000}^{xx}$ | $\hat{X}_{5001 5000}$ | $\Sigma_{5001 5000}^{xx}$ | $\Sigma_{5001 5000}^{yy}$ |
|-------|-----------------------|---------------------------|-----------------------|---------------------------|---------------------------|
| Value | 20.76 | $3.660 \cdot 10^{-3}$ | 20.76 | $1.366 \cdot 10^{-2}$ | $1.866 \cdot 10^{-2}$ |

Table 3: Final state of filter at observation 5000, where the outliers have been skipped.

Comparing table 3 to table 1, there appears to be no difference when it comes to the final state of the filter. This is likely because the outliers are relatively few and far away from the final state.

Figure 8 presents the ACF and PACF plots for the standardised residuals for the entire dataset. We see that there are some dynamics which are not captured by our model due to its simplicity. Hence, the assumptions of the iid residuals of the model are violated.

5 Optimizing the variances

Variance is a non-negative quantity, so naturally, both system and observation noise variance are lower bounded by zero.

Now, let's see if we can derive some initial estimates of the variances of the system and the observation noise.

We will begin with the observation noise variance. First, notice that salinity is measured in steps of size 0.01. Therefore, let's assume that the observation of salinity is just a rounding to the nearest 0.01, and that this is the *only* variance in the observation. Formally, this means that,

$$\begin{aligned}
\mathbf{Y}_t &= \text{round}(\mathbf{X}_t) \\
&= \mathbf{X}_t + \text{round}(\mathbf{X}_t) - \mathbf{X}_t \\
&= \mathbf{X}_t + \varepsilon_{2,t},
\end{aligned}$$

where $\varepsilon_{2,t} = \text{round}(\mathbf{X}_t) - \mathbf{X}_t$, and $\text{round}(\cdot)$ maps a number to the nearest element in $\mathbb{Z}10^{-2}$ (i.e. to nearest 0.01) with round to even for ties.

Since salinity is a physical quantity on a larger scale than atomic, it is fair to assume that its distribution has a continuous density. Now, we realize that from the definition of continuity, we can split the salinity scale into bins (closed intervals), such that the difference between max and min are as small as we like, just by making the bins small enough. This means that we can approximate its density function arbitrarily well by a weighted sum of *uniform* distributions.

Now, we notice that the variation of salinity is many orders of magnitude greater than 0.01, and therefore, we crudely assume that \mathbf{X}_t can be approximated well by a weighted sum of uniform distributions over intervals of type

$$\left[10^{-2} \left(n - \frac{1}{2} \right), 10^{-2} \left(n + \frac{1}{2} \right) \right], \quad n \in \mathbb{Z}$$

This means that, given one observation, \mathbf{Y}_t , if we assume that the only observation error is introduced by rounding, we know that $\mathbf{X}_t \in [Y_t - 0.005, Y_t + 0.005]$,

and it is uniformly distributed within this interval. This corresponds to

$$\varepsilon_{2,t} \sim \text{Uniform}\left(\frac{-10^{-2}}{2}, \frac{10^{-2}}{2}\right),$$

which is known to have variance

$$\text{Var}(\varepsilon_{2,t}) = \frac{1}{12} (10^{-2})^2 \approx 8.333 \cdot 10^{-6}.$$

Thus, under the assumption that the only error in \mathbf{Y}_t is the one introduced by rounding, and that the precision is high enough for the piecewise uniform approximation to be good enough; the observation noise variance is approximately $8.333 \cdot 10^{-6}$.

Now, we'll estimate an initial guess of system noise variance. First recall that $\mathbf{Y}_t = \mathbf{X}_t + \varepsilon_{2,t}$, and since \mathbf{X}_t is a random walk, $\mathbf{X}_t = \mathbf{X}_{t-1} + \varepsilon_{1,t} \Rightarrow \varepsilon_{1,t} = \mathbf{X}_t - \mathbf{X}_{t-1}$

$$\begin{aligned} \mathbf{Y}_t - \mathbf{Y}_{t-1} &= \mathbf{X}_t - \mathbf{X}_{t-1} + \varepsilon_{2,t} - \varepsilon_{2,t-1} \\ &= \varepsilon_{1,t} + \varepsilon_{2,t} - \varepsilon_{2,t-1} \Rightarrow \\ \varepsilon_{1,t} &= \mathbf{Y}_t - \mathbf{Y}_{t-1} - \varepsilon_{2,t} + \varepsilon_{2,t-1} \end{aligned}$$

Since we just saw that $\text{Var}(\varepsilon_{2,t}) \approx \frac{10^{-4}}{12}$, we get

$$\begin{aligned} \text{Var}(\varepsilon_{1,t}) &= \text{Var}(\mathbf{Y}_t - \mathbf{Y}_{t-1}) + \text{Var}(\varepsilon_{2,t}) + \text{Var}(\varepsilon_{2,t-1}) \\ &\approx \text{Var}(Y_t - Y_{t-1}) + \frac{10^{-4}}{6} \end{aligned}$$

There are no missing values in the first 800 observations, so we use these to estimate $\text{Var}(Y_t - Y_{t-1}) \approx 1.88891 \cdot 10^{-3}$. Noticing that this is multiple orders of magnitude greater than $2\text{Var}(\varepsilon_{2,t})$, we get

$$\begin{aligned} \text{Var}(\varepsilon_{1,t}) &\approx 1.88891 \cdot 10^{-3} + 8.333 \cdot 10^{-6} \\ &\approx 1.88891 \cdot 10^{-3} \end{aligned}$$

Thus our initial guess of system noise variance is $1.88891 \cdot 10^{-3}$.

Now, the maximum likelihood estimates of the two variances are obtained by minimizing [1, Eq. 10.149] according to the procedure described in [1, Chap. 10.6] with the log-transformed system and observation noise variance as θ . The minimization is performed via the R-function "optim" using the "SANN" method² with initial parameters $\log 0.01$ for both variances³. The results are found in table 4, and as we see, our initial guesses are actually very, very close to the ML-estimates. This indicates that the vast majority of the error stems from the model being wrong. This makes good sense, as it is a very

²The other methods seemed to push the observation noise variance to zero which is likely wrong because of the rounding.

³Our lecturer suggested to use/test with these values.

simple model. Furthermore, it indicates that the sensor is actually very precise, and the error is dominated by the rounding.

As mentioned earlier, we did experience some optimizers setting $\sigma_2 = 0$, which doesn't seem to make sense in light of the above contemplations about the variances. So let's elaborate a bit on what may have happened.

First, note that since our model is a random walk, it can be rewritten in the form:

$$\mathbf{Y}_t = \sum_{i=1}^t \varepsilon_{1,i} + \varepsilon_{2,t}$$

This means that

$$\begin{aligned} Var(\mathbf{Y}_t) &= Var\left(\sum_{i=1}^t \varepsilon_{1,i} + \varepsilon_{2,t}\right) \\ &= t\sigma_1 + \sigma_2 \end{aligned}$$

Therefore, since σ_2 is already much, much lower than σ_1 , the variance of \mathbf{Y}_t is by far dominated by the variance of the state. Thus, the optimizer may overfit to the data and just set it to zero while compensating with σ_1 - even though, clearly this does not give the same model.

| Parameter | ML | Guess |
|------------------|-----------------------|-----------------------|
| $\hat{\sigma}_1$ | $1.885 \cdot 10^{-3}$ | $1.889 \cdot 10^{-3}$ |
| $\hat{\sigma}_2$ | $6.765 \cdot 10^{-6}$ | $8.333 \cdot 10^{-6}$ |

Table 4: ML-estimates and initial guesses of system and observation noise variance based upon first 800 observations.




Figure 9: Predictions of the Kalman Filter with the skipping outliers mechanism and ML-estimated parameters, along with the prediction intervals in the range [800, 950]. The green lines are there to help orient readers in the zoomed in region.

| State | $\hat{X}_{5000 5000}$ | $\Sigma_{5000 5000}^{xx}$ | $\hat{X}_{5001 5000}$ | $\Sigma_{5001 5000}^{xx}$ | $\Sigma_{5001 5000}^{yy}$ |
|-------|-----------------------|---------------------------|-----------------------|---------------------------|---------------------------|
| Value | 20.77 | $6.741 \cdot 10^{-6}$ | 20.77 | $1.891 \cdot 10^{-3}$ | $1.898 \cdot 10^{-3}$ |

Table 5: Final state of filter at observation 5000. The previously found outliers have been skipped.

On table 5 it can be seen that compared to table 3, the variance of the reconstruction and prediction is orders of magnitude lower. The prediction, however, remains mostly the same - only difference is the last few decimals.

6 Model for dissolved oxygen

To build the model for DO_t , we start by modeling it in continuous time with a differential equation. Then, to get a discrete linear model, we use the first-order Taylor expansion, which is equivalent to the Euler's method.

The differential equation is the following:

$$\frac{d}{dt}DO(t) = \text{Production}(t) - \text{Consumption}(t),$$

where obviously $\text{Production}(t)$ is the rate of oxygen production at time t and $\text{Consumption}(t)$ is the oxygen consumption at time t .

The oxygen production (the primary production) is known to be a linear function of the sun intensity, $I(t)$, and the oxygen exchange with the atmosphere.

Before specifying this linear model, we note that the rate of oxygen exchange with the atmosphere, is known to behave such that $DO(t)$ approaches $DOsat(t)$. We assume that the rate is proportional to the negative signed difference, i.e.

$$\text{Exchange}(t) \propto DOsat(t) - DO(t)$$

Now we use this to get

$$\begin{aligned} \text{Production}(t) &= \alpha I(t) + \beta \text{Exchange}(t) \\ &= \alpha I(t) + \beta (DOsat(t) - DO(t)) \end{aligned}$$

where we have included the proportion constant in the parameter β . The oxygen consumption is just the respiration, $R(t)$, so now we can write the final differential equation:

$$\begin{aligned} \frac{d}{dt} DO(t) &= \text{Production}(t) - \text{Consumption}(t) \\ &= \underbrace{\alpha I(t) + \beta (DOsat(t) - DO(t))}_{\text{Production}} - \underbrace{R(t)}_{\text{Consumption}} \end{aligned}$$

Now we linearize using Taylor expansion, and discretize by setting $\Delta t = 1$.

$$\begin{aligned} DO(t + \Delta t) &= DO(t) + \Delta t \left(\frac{d}{dt} DO(t) \right) + \xi \Rightarrow \\ DO_{t+1} &= DO_t + \alpha I_t + \beta (DOsat_t - DO_t) - R_t + \xi \\ &= DO_t(1 - \beta) + \alpha I_t + \beta DOsat_t - R_t + \xi, \end{aligned}$$

where ξ is the error term. Note that the use of ξ is a bit sloppy, as it will later just be included in the noise term of the model. We include respiration as a random walk. That is, we have

$$\begin{cases} DO_t = DO_{t-1}(1 - \beta) - R_{t-1} + \beta DOsat_{t-1} + \alpha I_{t-1} + \epsilon_{(1,1),t} \\ R_t = R_{t-1} + \epsilon_{(1,2),t} \end{cases}$$

Equivalently, we can formulate it as the following state-space model *with* model error, $\epsilon_{1,t} := [\epsilon_{(1,1),t}, \epsilon_{(1,2),t}]^T \in \mathbb{R}^2$ and measurement error $\epsilon_{2,t} \in \mathbb{R}$:

$$\begin{aligned} \underbrace{\begin{bmatrix} DO_t \\ R_t \end{bmatrix}}_{\mathbf{X}_t} &= \underbrace{\begin{bmatrix} 1 - \beta & -1 \\ 0 & 1 \end{bmatrix}}_{\mathbf{A}} \underbrace{\begin{bmatrix} DO_{t-1} \\ R_{t-1} \end{bmatrix}}_{\mathbf{X}_{t-1}} + \underbrace{\begin{bmatrix} \beta & \alpha \\ 0 & 0 \end{bmatrix}}_{\mathbf{B}} \underbrace{\begin{bmatrix} DOsat_{t-1} \\ I_{t-1} \end{bmatrix}}_{\mathbf{u}_{t-1}} + \epsilon_{1,t} \\ \mathbf{Y}_t &= \underbrace{\begin{bmatrix} 1 & 0 \end{bmatrix}}_{\mathbf{C}} \underbrace{\begin{bmatrix} DO_t \\ R_t \end{bmatrix}}_{\mathbf{X}_t} + \epsilon_{2,t} \end{aligned}$$

References

- [1] Henrik Madsen. *Time Series Analysis*. Oct. 2008. ISBN: 978-1-4200-5967-0. DOI: 10.1201/9781420059687.

7 Code

7.1 A

```
1 ———
2 title: Assignment 4
3 author: sarphiv
4 output:
5   bookdown::html_document2: default
6   bookdown::pdf_document2: default
7 ———
8 ““{r include = FALSE}
9 # Disable inclusion of code chunks in the output
10 knitr::opts_chunk$set(echo = FALSE)
11 library("stats")
12 library("dplyr")
13 library("r2r")
14
15 set.seed(1337)
16
17
18 # Function to reset the default parameters
19 par_defaults <- par(no.readonly = TRUE)
20 par_reset <- function() {
21   par(par_defaults)
22 }
23 ““
24
25 ““{R}
26 # Load the data
27 data <- read.csv("A4_Kulhuse.csv", header = TRUE, sep =
    ",")
28 oxygen <- data$ODO
29 salinity <- data$Sal
30 time <- as.POSIXct(dplyr::pull(data, DateTime))
31 time_rel <- as.numeric(as.POSIXct(dplyr::pull(data,
    DateTime)))
32 time_rel <- (time_rel - time_rel[1])
33 time_delta <- mean(diff(time_rel), na.rm = T)
34 time_rel <- time_rel / time_delta
```

```

35
36 # Length
37 n <- length(time_rel)
38 ''
39
40 #
41 ''{R}
42 par(mfrow = c(2, 2))
43 plot(time, oxygen, type='l')
44 plot(time, salinity, type='l')
45 plot(oxygen, salinity)
46 boxplot(oxygen, salinity)
47 ''
48
49 #
50 Example 10.6 may be relevant for further detail.
51
52 Assuming stationary dynamics.
53 $$
54  $X_t = X_{t-1} + \epsilon_t$ 
55 $$
56
57 Must find variance of  $\epsilon_t \sim \mathcal{N}(0, \sigma^2)$ .
58
59 It is assumed that the salinity is measured directly,
    but with independent measurement noise.
60 $$
61  $Y_t = X_t + \eta_t$ 
62 $$
63
64 According to section 10.1, the matrices are given as
65 $$
66  $A = I$ 
67  $B = 0$ 
68  $C = I$ 
69 $$
70
71
72 #
73 ''{R}
74 kalman <- function(A, C, X.var, Y.var, X0.pred,
    X0.pred.var, Y, n.ahead=1, B=NULL, u=NULL,
    outlier.sd=NULL, outlier.idx=NULL, verbose=NULL) {
75   # NOTE: Implementation is according to theorem 10.2
76   # Assuming stationary dynamics.

```

```

77
78 # Initialize state and variance variables
79 n <- length(Y) + n.ahead
80
81 X.dim <- dim(X0.pred)[1]
82 Y.dim <- dim(C)[1]
83
84 X.recon <- lapply(1:n, function(x) rep(NA, X.dim))
85 X.recon.var <- lapply(1:n, function(x) matrix(NA, nrow
  = X.dim, ncol = X.dim))
86
87 X.pred <- lapply(1:n, function(x) rep(NA, X.dim))
88 X.pred.var <- lapply(1:n, function(x) matrix(NA, nrow
  = X.dim, ncol=X.dim))
89
90 Y.pred.var <- lapply(1:n, function(x) matrix(NA, nrow
  = Y.dim, ncol = Y.dim))
91
92
93 # Initialize optional arguments
94 if (is.null(u)) {
95   u <- matrix(0, nrow = 1, ncol = n)
96 }
97 if (is.null(B)) {
98   B <- matrix(0, nrow = X.dim, ncol = dim(u)[1])
99 }
100
101 # Outlier handling
102 if (is.null(outlier.idx)) {
103   outlier.idx <- hashset()
104 }
105
106 outlier.idx.detected <- hashset()
107
108
109 # Helper functions
110 recon.calc <- function(i) {
111   # If missing observation, set reconstruction to
  prediction
112   if (i > length(Y) || any(is.na(Y[[i]]))) {
113     X.recon[[i]] <- X.pred[[i]]
114     X.recon.var[[i]] <- X.pred.var[[i]]
115
116     if ("missing" %in% verbose) {
117       cat("Missing:", i, "\n")
118     }

```

```

119
120     return()
121 }
122
123
124     # Mahalanobis distance for outlier detection
125     Y.error <- abs(Y[[i]] - C %*% X.pred[[i]])
126     Y.error.sd <- sqrt((t(Y.error) %*%
solve(Y.pred.var[[i]]) %*% Y.error) / Y.dim)
127
128     # Outlier and outlier_sd is set
129     outlier.sd_detected <- !is.null(outlier.sd) &&
(Y.error.sd > outlier.sd)
130     # Outlier index is marked
131     outlier.marked <- outlier.idx[[i]]
132     # If outlier, set reconstruction to prediction
133     if (outlier.sd_detected || outlier.marked) {
134         X.recon[[i]] <-< X.pred[[i]]
135         X.recon.var[[i]] <-< X.pred.var[[i]]
136
137         if (outlier.sd_detected) {
138             outlier.idx.detected[[i]] <-< TRUE
139         }
140
141         if ("outlier" %in% verbose) {
142             cat("Outlier:", i, "Sd:", Y.error.sd, "\n")
143         }
144     }
145     # Else, reconstruct
146     else {
147         K <-< X.pred.var[[i]] %*% t(C) %*%
solve(Y.pred.var[[i]])
148
149         X.recon[[i]] <-< X.pred[[i]] + K %*% (Y[[i]] - C
%*% X.pred[[i]])
150         X.recon.var[[i]] <-< X.pred.var[[i]] - K %*% C %*%
X.pred.var[[i]]
151     }
152
153 }
154
155 pred.calc <- function(i) {
156     X.pred[[i]] <-< A %*% X.recon[[i-1]] + B %*% u[[i-1]]
157     X.pred.var[[i]] <-< A %*% X.recon.var[[i-1]] %*%
t(A) + X.var
158     Y.pred.var[[i]] <-< C %*% X.pred.var[[i]] %*% t(C) +

```



```

    Y.var
159 }
160
161
162 # NOTE: Initial predictions set according to theorem
      10.2.
163 # Not setting reconstruction directly as in lecture
      notes,
164 # as it would not incorporate dynamics and the first
      observation.
165 X.pred[[1]] <- X0.pred
166 X.pred.var[[1]] <- X0.pred.var
167
168 Y.pred.var[[1]] <- C %*% X.pred.var[[1]] %*% t(C) +
      Y.var
169
170 # Initial reconstruction
171 recon.calc(1)
172
173
174 # Kalman filter iterative implementation of recursion
175 for (i in 2:n) {
176   pred.calc(i)
177   recon.calc(i)
178 }
179
180
181 # Return reconstructed and predicted values
182 return(list(
183   X.recon      = X.recon ,
184   X.recon.var  = X.recon.var ,
185   X.pred       = X.pred ,
186   X.pred.var   = X.pred.var ,
187   Y.pred.var   = Y.pred.var ,
188   outlier.idx  = outlier.idx.detected
189 ))
190 }
191
192
193 results.pure <- kalman(
194   A=matrix(1, nrow = 1, ncol = 1) ,
195   C=matrix(1, nrow = 1, ncol = 1) ,
196   X.var=matrix(0.01, nrow = 1, ncol = 1) ,
197   Y.var=matrix(0.005, nrow = 1, ncol = 1) ,
198   X0.pred=matrix(salinity[1], nrow = 1, ncol = 1) ,
199   X0.pred.var=matrix(var(salinity , na.rm = T), nrow = 1 ,

```

```

        ncol = 1),
200   Y = as.list(salinity)
201 )
202
203 ““
204
205 ““{R}
206 par_reset()
207 plot(time, salinity, type = "l")
208 lines(time, results.pure$X.pred[-(n+1)], col = "red")
209 ““
210
211 ““{R}
212 par_reset()
213 plot(time[800:950], salinity[800:950], type = "l")
214 lines(time[800:950], results.pure$X.pred[800:950], col =
      "red")
215 ““
216
217 ““{R}
218 par_reset()
219 plot(time[800:950], (salinity[800:950] -
      vapply(results.pure$X.pred[800:950], function(x) x[1,
      1], c(1))) /
      sqrt(vapply(results.pure$X.pred.var[800:950],
      function(x) x[1, 1], c(1))), type = "p")
220 ““
221
222 ““{R}
223 print(results.pure$X.recon[[5000]])
224 print(results.pure$X.recon.var[[5000]])
225 print(results.pure$X.pred[[5001]])
226 print(results.pure$X.pred.var[[5001]])
227 print(results.pure$Y.pred.var[[5001]])
228 ““
229
230 #
231 ““{R}
232 results.skipped <- kalman(
233   A=matrix(1, nrow = 1, ncol = 1),
234   C=matrix(1, nrow = 1, ncol = 1),
235   X.var=matrix(0.01, nrow = 1, ncol = 1),
236   Y.var=matrix(0.005, nrow = 1, ncol = 1),
237   X0.pred=matrix(salinity[1], nrow = 1, ncol = 1),
238   X0.pred.var=matrix(var(salinity, na.rm = T), nrow = 1,
      ncol = 1),

```

```

239 Y=as.list(salinity),
240 outlier.sd=6,
241 verbose=c("outlier")
242 )
243
244 outlier.idx <-
      sort(unlist(keys(results.skipped$outlier.idx)))
245 ""
246
247 ""{R}
248 par_reset()
249 plot(time, salinity, type = "l")
250 for (i in outlier.idx) {
251   abline(v = time[i], col = "blue")
252 }
253 abline(v=time[is.na(salinity)], col="#00ff002c")
254 lines(time, results.skipped$X.pred[-(n + 1)], col =
      "red")
255 ""
256
257 ""{R}
258 par_reset()
259 plot(time[-outlier.idx], salinity[-outlier.idx], type =
      "l")
260 for (i in outlier.idx) {
261   abline(v = time[i], col = "blue")
262 }
263 abline(v=time[is.na(salinity)], col="#00ff002c")
264 lines(time[-outlier.idx], results.skipped$X.pred[-(n +
      1)][-outlier.idx], col = "red")
265 ""
266
267 ""{R}
268 par_reset()
269 plot(time[800:950], salinity[800:950], type = "l")
270 lines(time[800:950], results.skipped$X.pred[800:950],
      col = "red")
271 ""
272
273 ""{R}
274 par_reset()
275 plot(time[800:950], (salinity[800:950] -
      vapply(results.skipped$X.pred[800:950], function(x)
        x[1, 1], c(1))) /
      sqrt(vapply(results.skipped$X.pred.var[800:950],
        function(x) x[1, 1], c(1))), type = "p")

```

```

276 ""
277
278 ""{R}
279 print(results.skipped$X.recon[[5000]])
280 print(results.skipped$X.recon.var[[5000]])
281 print(results.skipped$X.pred[[5001]])
282 print(results.skipped$X.pred.var[[5001]])
283 print(results.skipped$Y.pred.var[[5001]])
284 ""
285
286 #
287
288 ## Non-parametric bootstrap of observation rounding
289 ""{R}
290 obs <- runif(1000000)
291 obs.round <- round(obs, 2)
292 obs.error <- obs.round - obs
293
294 par(mfrow = c(1, 2))
295 hist(obs, breaks = 100, main = "Histogram of original
    decimals", xlab = "Decimal")
296 hist(obs.error, breaks = 100, main = "Histogram of
    error", xlab = "Error")
297
298 print(var(obs.error))
299 print(0.01^2 / 12)
300
301 print(var(salinity[1:800]))
302 print(var(diff(salinity[1:800])))
303 ""
304
305 Seems like the error is uniformly distributed between
    -0.005 and 0.005.
306 The lower bound is therefore 0.
307
308 A crude estimate of the system noise order of magnitude
    is also given.
309
310 ""{R}
311 log.likelihood <- function(vars) {
312   # WARN: Assuming no observations are missing
313   A <- matrix(1, nrow = 1, ncol = 1)
314   C <- matrix(1, nrow = 1, ncol = 1)
315   Y.var <- matrix(exp(vars[2]), nrow = 1, ncol = 1)
316
317   results <- kalman(

```

```

318   A = A,
319   C = C,
320   X.var = matrix(exp(vars[1]), nrow = 1, ncol = 1),
321   Y.var = Y.var,
322   X0.pred = matrix(salinity[1], nrow = 1, ncol = 1),
323   X0.pred.var = matrix(var(salinity, na.rm = T), nrow
= 1, ncol = 1),
324   Y = as.list(salinity[1:800]),
325   n.ahead = 0,
326   outlier.idx = results.skipped$outlier.idx,
327   verbose = c("outlier", "missing")
328 )
329
330 # Implementation according to book
331 # R <- lapply(results$X.pred.var, function(x) C %*% x
%*% t(C) + Y.var)
332 # # NOTE: Using 801 observations, since filtering with
1:800, but log likelihood with 2:801
333 # error <- lapply(1:800, function(i) salinity[i+1] - C
%*% A %*% results$X.recon[[i]])
334 # # NOTE: Swapped sign to maximize via optim, which
minimizes
335 # return(sum(vapply(1:799, function(i)
log(det(R[[i+1]])) + t(error[[i]]) %*%
solve(R[[i+1]]) %*% error[[i]], c(1))))
336
337
338 # Implementation according to slides
339 # error <- lapply(1:799, function(i) salinity[i+1] - C
%*% A %*% results$X.recon[[i]])
340 # NOTE: Swapped sign to maximize via optim, which
minimizes
341 # return(sum(vapply(2:800, function(i)
log(det(results$Y.pred.var[[i]])) + t(error[[i-1]])
%*% solve(results$Y.pred.var[[i]]) %*% error[[i-1]],
c(1))))
342
343
344 # Implementation simplified
345 error <- lapply(1:800, function(i) salinity[i] - C %*%
results$X.pred[[i]])
346 # NOTE: Swapped sign to maximize via optim, which
minimizes
347 return(sum(vapply(1:800, function(i)
log(det(results$Y.pred.var[[i]])) + t(error[[i]]) %*%
solve(results$Y.pred.var[[i]]) %*% error[[i]], c(1))))

```

```

348 }
349
350 # print(nlminb(c(log(0.01), log(0.01)), log.likelihood))
351 # print(nlm(log.likelihood, c(log(0.01), log(0.01))))
352 # print(optim(c(log(0.01), log(0.01)), log.likelihood,
353             control = list(trace = 1, REPORT = 1)))
353 optim.results <- optim(c(log(0.01), log(0.01)),
354                       log.likelihood, method = "SANN", control = list(maxit
355                             = 1000, tmax = 10, trace = 1, REPORT = 10))
354
355 print(optim.results$par)
356 print(exp(optim.results$par))
357 print(optim.results$value)
358 ""
359
360
361 ""{R}
362 results.optimal <- kalman(
363   A=matrix(1, nrow = 1, ncol = 1),
364   C=matrix(1, nrow = 1, ncol = 1),
365   X.var=matrix(exp(optim.results$par[1]), nrow = 1, ncol
366               = 1),
367   Y.var=matrix(exp(optim.results$par[2]), nrow = 1, ncol
368               = 1),
369   X0.pred=matrix(salinity[1], nrow = 1, ncol = 1),
370   X0.pred.var=matrix(var(salinity, na.rm = T), nrow = 1,
371                     ncol = 1),
372   Y=as.list(salinity),
373   outlier.idx = results.skipped$outlier.idx,
374   verbose = c("outlier")
375 )
376 ""
377 ""{R}
378 par_reset()
379 plot(time, salinity, type = "l")
380 lines(time, results.optimal$X.pred[-(n+1)], col = "red")
381 ""
382 ""{R}
383 par_reset()
384 plot(time[800:950], salinity[800:950], type = "l")
385 lines(time[800:950], results.optimal$X.pred[800:950],
386       col = "red")
387 ""

```

```

387 ““{R}
388 par_reset()
389 plot(time[800:950], (salinity[800:950] -
      vapply(results$optimal$X.pred[800:950], function(x)
        x[1, 1], c(1))) /
      sqrt(vapply(results$optimal$X.pred.var[800:950],
        function(x) x[1, 1], c(1))), type = "p")
390 ““
391
392 ““{R}
393 print(results$optimal$X.recon[[5000]])
394 print(results$optimal$X.recon.var[[5000]])
395 print(results$optimal$X.pred[[5001]])
396 print(results$optimal$X.pred.var[[5001]])
397 print(results$optimal$Y.pred.var[[5001]])
398 ““
399
400
401 #
402 ““{R}
403
404 ““

```